

Chapter 11 – Neural Nets

**Machine Learning for Business
Analytics in R (2nd ed)**

**Shmueli, Bruce, Gedeck, Yahav, &
Patel**


Basic Idea

- Combine input information in a complex & flexible neural net “model”
- Model “coefficients” are continually tweaked in an iterative process
- The network’s interim performance in classification and prediction informs successive tweaks

Network Structure

- Multiple layers
 - Input layer (raw observations)
 - Hidden layers
 - Output layer
- Nodes
- Weights (like coefficients, subject to iterative adjustment)
- Bias values (also subject to iterative adjustment)

Schematic Diagram




Tiny Example

Predict consumer opinion of cheese product based
on fat and salt content

Obs.	Fat Score	Salt Score	Opinion
1	0.2	0.9	like
2	0.1	0.1	dislike
3	0.2	0.4	dislike
4	0.2	0.5	dislike ¹
5	0.4	0.5	like
6	0.3	0.8	like

Example - Using fat & salt content to predict consumer acceptance of cheese



Circles are nodes, w_{ij} on arrows are weights, and Θ_j are node bias values

Tiny Example - Data

<i>Obs.</i>	<i>Fat Score</i>	<i>Salt Score</i>	<i>Acceptance</i>
1	0.2	0.9	1
2	0.1	0.1	0
3	0.2	0.4	0
4	0.2	0.5	0 4
5	0.4	0.5	1
6	0.3	0.8	1

Moving Through the Network

The Input Layer

- For input layer, input = output
- E.g., for record #1:
 - Fat input = output = 0.2
 - Salt input = output = 0.9
- Output of input layer = input into hidden layer

The Hidden Layer

- In this example, it has 3 nodes
- Each node receives as input the output of all input nodes
- Output of each hidden node is some function of the weighted sum of inputs

$$output_j = g(\Theta_j + \sum_{i=1}^p w_{ij} x_i)$$

The Weights


- The weights θ (theta) and w are typically initialized to random values in the range -0.05 to +0.05
- Equivalent to a model with random prediction (in other words, no predictive value)
- These initial weights are used in the first round of training

Output of Node 3 if g is a Logistic Function

$$output_j = g(\theta_j + \sum_{i=1}^p w_{ij} x_i)$$

$$output_3 = \frac{1}{1 + e^{-[-0.3 + (0.05)(0.2) + (0.01)(0.9)]}} = 0.43$$

Initial Pass of the Network



Node outputs for the first record in tiny example, and logistic function. Final quasi-propensities (0.506 for like and 0.481 for dislike) are normalized to 0.51 and 0.49 so they add to 1 (this operation is called *softmax*).

Output Layer

- The output of the last hidden layer becomes input for the output layer
- Uses same function as above, i.e. a function g of the weighted average

$$\text{Output}_{N6} = \frac{1}{1 + e^{-[-0.04 + (-0.02)(0.43) + (-0.03)(0.51) + \cancel{(0.015)(0.52)}}]} = 0.481$$
$$\text{Output}_{N7} = \frac{1}{1 + e^{-[-0.015 + (0.01)(0.43) + (0.05)(0.51) + (0.015)(0.52)}}} = 0.506.$$

Mapping the output to a classification

- Output = 0.506 for “like” and 0.481 for “dislike”
- Softmax normalization to 0.51 for “like” and 0.49 for “dislike”
- So classification, at this early stage, is “like”

Relation to Linear Regression

A net with a single output node and no hidden layers, where g is the identity function, takes the same form as a linear regression model

$$\hat{y} = \Theta + \sum_{i=1}^p w_i x_i$$

Training the Model

Preprocessing Steps

- Scale variables to 0-1
- Categorical variables
- If equidistant categories, map to equidistant interval points in 0-1 range
- Otherwise, create dummy variables
- Transform (e.g., log) skewed variables

Initial Roundtrip Through Network

- Goal: Find weights that yield best predictions
- The process we described above is repeated for all records
- At each record compare prediction to actual
- Difference is the error for the output node
- Error is propagated back and distributed to all the hidden nodes and used to update their weights

Back Propagation (“back-prop”)

- Output from output node k \hat{y}_k
- Error associated with that node:

$$err_k = \hat{y}_k(1 - \hat{y}_k)(y_k - \hat{y}_k)$$

Note: this is like ordinary error, multiplied by a correction factor

Error is Used to Update Weights

$$\theta_j^{new} = \theta_j^{old} + l(err_j)$$

$$w_j^{new} = w_j^{old} + l(err_j)$$

l = constant between 0 and 1, reflects the “learning rate” or “weight decay parameter”

Why It Works

- Big errors lead to big changes in weights
- Small errors leave weights relatively unchanged
- Over thousands of updates, a given weight keeps changing until the error associated with that weight is negligible, at which point weights change little

R Code for the Tiny Example

neuralnet (used here)

nnet (does not support multilayer networks)

Code for the tiny example:

```
nn <- neuralnet(Like + Dislike ~ Salt + Fat, data = df,  
linear.output = F, hidden = 3)  
  
# display weights  
nn$weights  
# display predictions  
prediction(nn)  
# plot network  
plot(nn, rep="best")
```




Indicates 1 hidden layer with 3 nodes, the syntax `hidden = 3,4` would mean 2 layers, 3 nodes in the first, 4 in the second

Predictions

```
$rep1
```

	Salt	Fat	Like	Dislike
1	0.1	0.1	0.0002415536	0.99965512
2	0.4	0.2	0.0344215787	0.96556788
3	0.5	0.2	0.1248666748	0.87816828
4	0.9	0.2	0.9349452648	0.07022732
5	0.8	0.3	0.9591361793	0.04505631
6	0.5	0.4	0.8841904620	0.12672438

Network With Final Weights



Confusion Matrix

(use predict from caret library)

```
library(caret)
predict <- compute(nn, data.frame(df$Salt, df$Fat))
predicted.class=apply(predict$net.result,1,which.max)-1
confusionMatrix(factor(ifelse(predicted.class=="1",
  "dislike", "like")), factor(df$Acceptance))
```

Output

```
> confusionMatrix(as.factor(ifelse(predicted.class=="1",
  "dislike", "like")),
+ df$Acceptance)
```

Confusion Matrix

		Reference	
		dislike	like
Prediction	dislike	3	0
like	0	3	

Common Criteria to Stop the Updating

- When weights change very little from one iteration to the next
- When the misclassification rate reaches a required threshold
- When a limit on runs is reached

Avoiding Overfitting

With sufficient iterations, neural net can easily overfit the data

To avoid overfitting:

- Track error in validation data
- Limit iterations
- Limit complexity of network

User Inputs

Specify Network Architecture

Number of hidden layers

- Most popular – one hidden layer

Number of nodes in hidden layer(s)

- More nodes capture complexity, but increase chances of overfit

Number of output nodes

- For classification with m classes, use m or $m-1$ nodes
- For numerical prediction use one

Network Architecture, cont.

“Learning Rate”

- Low values “downweight” the new information from errors at each iteration
- This slows learning, but reduces tendency to overfit to local structure

“Momentum”

- High values keep weights changing in same direction as previous iteration
- Likewise, this helps avoid overfitting to local structure, but also slows learning

Arguments in neuralnet

- `hidden`: a vector specifying the number of nodes per layer (thus specifying both the size and number of layers)
- `learningrate`: value between 0 and 1

Advantages

- Good predictive ability
- Can capture complex relationships
- No need to specify a model

Disadvantages

- Considered a “black box” prediction machine, with no insight into relationships between predictors and outcome
- No variable-selection mechanism, so you have to exercise care in selecting variables
- Heavy computational requirements if there are many variables (additional variables dramatically increase the number of weights to calculate)

Deep Learning

- The statistical and machine learning models in this book - including standard neural nets - work where you have informative predictors (purchase information, bank account information, # of rooms in a house, etc.)
- In rapidly-growing applications of voice and image recognition, you have high numbers of “low-level” granular predictors - pixel values, wave amplitudes, uninformative at this low level

Deep Learning

The most active application area for neural nets


- In image recognition, pixel values are predictors, and there might be 100,000+ predictors – big data! (voice recognition similar)
- Deep neural nets with many layers (“neural nets on steroids”) have facilitated revolutionary breakthroughs in image/voice recognition, and in artificial intelligence (AI)
- Key is the ability to self-learn features (“unsupervised”)
- For example, clustering could separate the pixels in this 1” by 1” football field image into the “green field” and “yard marker” areas without knowing that those concepts exist
- From there, the concept of a boundary, or “edge” emerges
- Successive stages move from identification of local, simple features to more global & complex features



Convolutional Neural Net

example in image recognition

- A popular deep learning implementation is a convolutional neural net (CNN)
- Need to aggregate predictors (pixels)
- Rather than have weights for each pixel, group pixels together and apply the same operation: “convolution”
- Common aggregation is a 3×3 pixel area, for example the small area around this man’s lower chin



Enlargement
of area

25	200	25
25	225	25
25	225	25

Pixel values (higher
number = darker)

Apply the convolution

Convolution operation is “multiply the pixel matrix by the filter matrix” then sum

0	1	0
0	1	0
0	1	0

x

25	200	25
25	225	25
25	225	25

$$\begin{aligned} 0*25 + 1*200 + 0*25 + \\ 0*25 + 1*225 + 0*25 + \\ 0*25 + 1*225 + 0*25 \\ = 650 \end{aligned}$$

Filter matrix that is good at identifying center vertical lines (we will see why shortly)

Pixel values

Sum = 650; this is higher than for any other arrangement of the filter matrix, because pixel values are highest in central column


Continue the Convolution

- The filter matrix moves across the image, storing its result, yielding a smaller matrix whose values indicate the presence or absence of a vertical line.
- Similar filters can detect horizontal lines, curves, borders - *hyper-local features*
- Further convolutions can be applied to these local features
- Result: multi-dimensional matrix, or **tensor**, of higher-level features

The Learning Process

How does the net learn which convolutions to do?


- In supervised learning, the net retains those convolutions and features which are successful in labeling (tagging) images
- Note that the feature-learning process yields a reduced (simpler) set of features than the original set of pixel values



Unsupervised Learning

Autoencoding

- Deep learning nets can learn higher level features even when there are no labels to guide the process
- The net adds a process to take the high level features and generate an image
- The generated image is compared to the original image and the net retains the architecture that produces the best matches



Deep Learning

The most active application area for neural nets

- In image recognition, pixel values are predictors, and there might be 100,000+ predictors – big data! (voice recognition similar)
- Deep neural nets with many layers (“neural nets on steroids”) have facilitated revolutionary breakthroughs in image/voice recognition, and in artificial intelligence (AI)
- Key is the ability to self-learn features (“unsupervised”)
- For example, clustering could separate the pixels in this 1” by 1” football field image into the “green field” and “yard marker” areas without knowing that those concepts exist
- From there, the concept of a boundary, or “edge” emerges
- Successive stages move from identification of local, simple features to more global & complex features



Deep Learning in R

The TensorFlow and Keras implementations in R require a Python environment with both installed. Find up to date installation instructions at <https://tensorflow.rstudio.com>. Building deep learning models requires the following three packages:

reticulate Interface to ‘Python’ modules, classes, and functions.

tensorflow R interface to ‘TensorFlow’ (<https://www.tensorflow.org>), an open source software library for numerical computation using data flow graphs

keras R interface to ‘Keras’ (<https://keras.io>), a high-level neural networks API

Summary

- Neural networks can be used for classification and prediction
- Can capture a very flexible/complicated relationship between the outcome and a set of predictors
- The network “learns” and updates its model iteratively as more data are fed into it
- Major danger: overfitting
- Requires large amounts of data
- Good predictive performance, yet “black box” in nature
- Deep learning (very complex neural nets), is effective in image recognition and AI